

Regression

PHYS591000 2022.03.02

Outline


- Linear regression: Simply fit a line!
- Linear regression with multivariable and polynomials
- Problem of overfitting and regularization
- Non-linear? Ask the neighbors!
 - the k-Nearest Neighbors (kNN) Algorithm

Warming up

- Access control of the building will be granted soon.
- As usual, take 3 mins to introduce yourself to your teammate for this week!
 - “So you’re really staying in this class!”
 - “How’s your experience so far? What do you think we can do to make this class more enjoyable?”


Regression

- Start with the simplest case: Fit a line!
→ Linear regression
- Physics example: Predict energy of a particle using information of a calorimeter




Linear Regression

- Calorimeters are used to measure the energy of a particle
- The incoming particle interacts with materials in the calorimeter and produce a bunch of other particles. → “Shower”




Linear Regression

- Expect the energy of the particle depends on the properties of the shower it creates, e.g. number of particles in the shower




Linear Regression

- Goal: Predict the energy of a particle given the number of shower particles it creates.




Linear Regression

- Fit a line!




Linear Regression

- Make a prediction of y (**target**) given the value of x (**feature**) using a hypothesis $y = h(x) = \theta_0 + \theta_1 x$



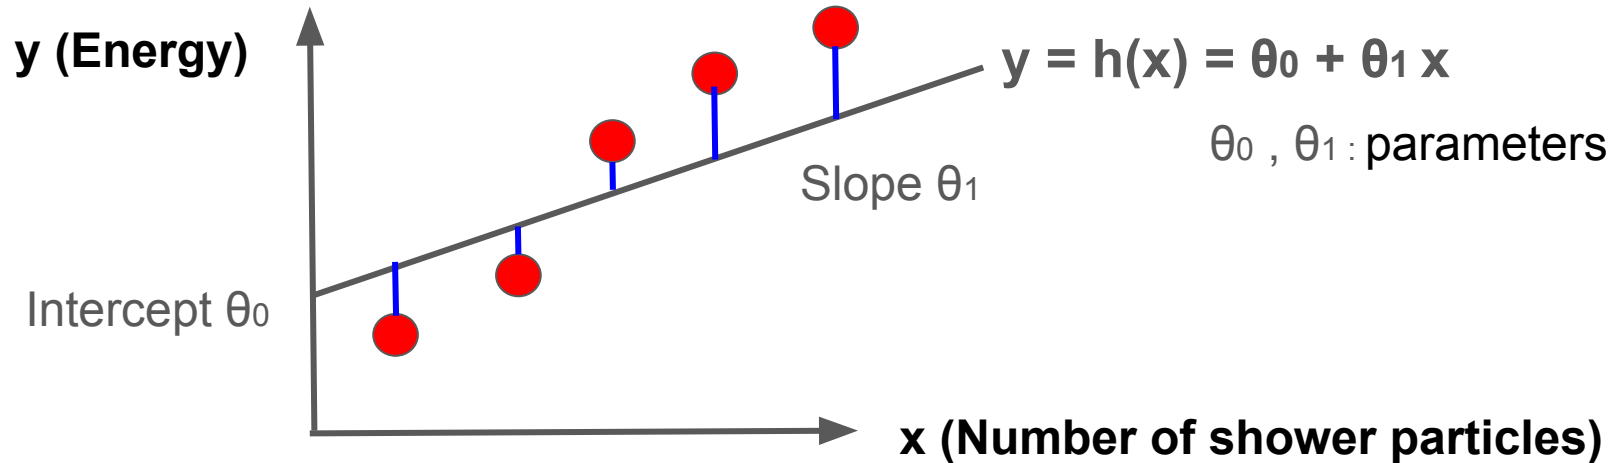
Linear Regression

- The line $y = h(x) = \theta_0 + \theta_1 x$ is the line that gives the *closest* predictions to all the points from the training sample



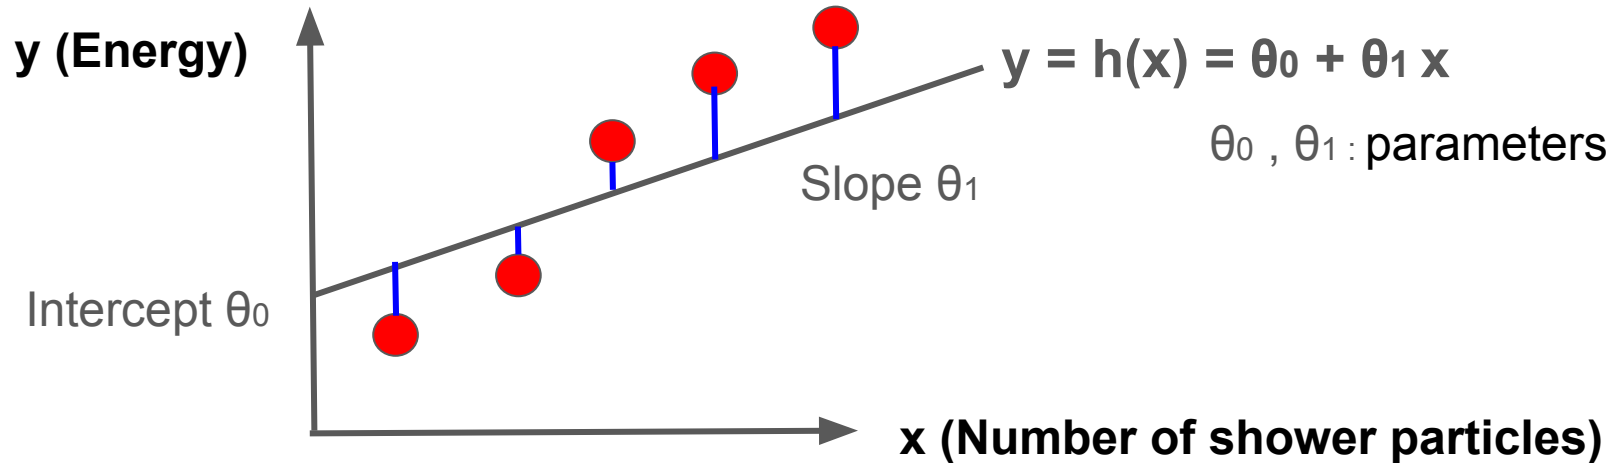
Linear Regression

- Find θ_0 and θ_1 that minimizes the sum of the squares $\sum_i (h(x_i) - y_i)^2$ from each (x_i, y_i) of the training sample \rightarrow method of least squares




Linear Regression

- $\sum_i (h(x_i) - y_i)^2$ is called the objective function
(or **cost** function, **loss** function)




Linear Regression

- Can extend to multivariable (more features) regression
→ Minimize $\sum_i (h(x_i, z_i) - y_i)^2$



Linear Regression

- What if we want to fit with a polynomial function?




Linear Regression


- What if we want to fit with a polynomial function?
→ Same as fitting with multiple features!

$$y = h(x) = \theta_0 + \theta_1 x + \theta_2 x^2$$

$$y = h(x, z) = \theta_0 + \theta_1 x + \theta_2 z$$




Linear Regression



$$y = h(x) = \theta_0$$

Underfitting



$$y = h(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$


Overfitting

Linear Regression

Training

Testing

$$y = h(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$




Overfitting: When there are too many parameters/features, the learned hypothesis may fit the training set ‘too well’ (sum of least square $\rightarrow 0$), but fail to generalize to new samples (fail to make good predictions for test samples).

Overfitting

- Options to avoid/fix overfitting:
 - Manually reduce the number of features/parameters
 - **Regularization:** Add a 'penalty' for the sizes of parameters


Regularization

- Example: overfitting from two points \rightarrow Sum of squares
$$\sum_i (h(x_i) - y_i)^2 = 0$$



Regularization


- The line fits training (red) samples perfectly but fails to generalize to test samples (blue)



Regularization

- Regularization: Minimize $\sum_i (h(x_i) - y_i)^2 + \lambda \theta_1^2$ λ : hyperparameter


Slope θ_1 (Parameter)



Regularization


- Regularization: Minimize $\sum_i (h(x_i) - y_i)^2 + \lambda \theta_1^2$ λ : hyperparameter

Slope θ_1 (Parameter)




Regularization

- Regularization: Minimize $\sum_i (h(x_i) - y_i)^2 + \lambda \theta_1^2$ λ : hyperparameter




Regularization

- Regularization: Minimize $\sum_i (h(x_i) - y_i)^2 + \lambda \theta_1^2$ λ : hyperparameter



Regularization

- Regularization: Minimize $\sum_i (h(x_i) - y_i)^2 + \lambda \theta_1^2$ λ : hyperparameter



Regularization

- Regularization: Minimize $\sum_i (h(x_i) - y_i)^2$ (with θ_0) + $\lambda \sum_i \theta_i^2$
(from $i=1$, i.e. without θ_0 ; we do not penalize the overall constant θ_0)
- λ : **hyperparameter** chosen with an independent **validation** sample
→ Then apply on another independent test sample to evaluate the performance
- Ridge regression: $\lambda \sum_i \theta_i^2$ (from $i=1$)
Lasso regression: $\lambda \sum_i |\theta_i|$ (from $i=1$)

Linear Regression: Datasets

- For both in-class exercise and lab this week we'll use the data from a neutrino experiment called OPERA.
- Let's turn to this week's in-class exercise to play with what we've learned.


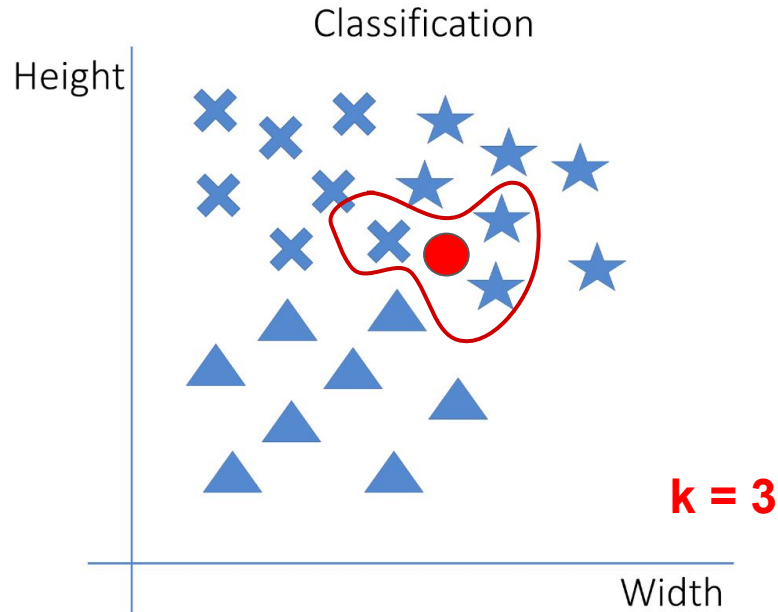
Non-linear models

- So far we've been working with linear models:
 - Classify two categories with a line (LDA)
 - Making continuous predictions (Regression) by fitting a line (or a curve)
- What if the distributions are not so linear? (As you've already seen for N_{\max} and iz_{\max} in the OPERA data)

The k-Nearest Neighbors (kNN) Algorithm


- A simple method is to ‘make it the same as its neighbors’:
 - Pick a positive integer k ($=1,2,3,\dots$)
 - For a test (unknown) point, find the k -nearest neighbors of (known/training) data
 - Classification: take votes from the neighbors and classify the test data as the same category which gets the most votes
 - Regression: take average target value of the neighbors

The k-Nearest Neighbors (kNN) Algorithm



The k-Nearest Neighbors (kNN) Algorithm

- kNN is intuitive and fast (no need to 'train' the model)
- The results depends on the choice of k



The k-Nearest Neighbors (kNN) Algorithm

- k too small: The prediction may be affected by noise/outliers
- k too large: Lose sensitivity to categories of small numbers
- Often need to try out a few different k, and validate the performance with test samples.

Lab for this week

- For the Lab this week, you'll keep studying regression from the task of predicting neutrino energy with OPERA data, and compare the performance of linear regressions to kNN!

Summary-I


- Linear regression: Fit the training sample with a line by minimizing the sum of squares $\sum_i (h(x_i) - y_i)^2$
- Linear regression can be generalized to cases with multivariable (more than one feature) or with polynomials.

Summary-II

- When there are too many features one may overfit the training sample
 - Consequence: Cannot generalize to test samples
- Regularization: Add a penalty for the size of the parameters to make the prediction less sensitive to the training sample
- E.g. Ridge regression: minimize $\sum_i (h(x_i) - y_i)^2 + \lambda \sum_i \theta_i^2$ (from $i=1$)
 - The hyperparameter λ can be tuned with validation samples

Addendum

- Overfitting may happen for classification too!
→ Sometimes called “overtraining”
- Similarly one can ‘regularize’ the classification algorithm. We’ll touch upon the relevant ideas later.



Source: Wikipedia