

Convolutional Neural Networks

PHYS591000 2022.04.13

Announcement

- It is possible that we need to go online (remote teaching) at some point this semester due to the pandemic.
- If so, we'll use Google Meet.
- Please check your email regularly for the rest of the semester for any sudden change of schedule.

Warming up

- As usual, take 3 mins to introduce yourself to your teammate for this week!
 - “How was your spring break?”
 - “Have you decided on the final project?”

Outline

- Convolution neural network (CNN) is particularly useful for tasks related to images. Why? How does it work?
- The example for today's lecture is taken from [this video](#) of the youtube channel 'StatQuest'. You are strongly encouraged to check out this extraordinary channel of statistical science and machine learning.

DNN from Week 06

- We'll build a 784-30-10 NN:

of biases:

0 (no biases for input)+

30 (hidden neurons)+

10 (output neurons).

of weights:


784×30 (input to hidden)+

30×10 (hidden to output).

23860 parameters in total



Lots of parameters to be optimized in the training!




DNN v.s. CNN for image classification

- Before we take every pixel as an independent input → each pixel has a weight and a bias for each neuron in the next layer.
 - When the NN becomes deeper and deeper (more and more hidden layers) the parameters will explode, and takes a lot of time for training.
- Need a different NN structure.

DNN v.s. CNN for image classification

- Another useful feature that is 'lost' in the DNN approach is the information on the *correlations* among pixels in an image:



Typically black pixels are near other black pixels.



There are certain patterns of 'edges' between black and white pixels (e.g. for a face).

CNN for Image Classification

- Task: Classify images of O and X
 - Each image is 6x6 pixel



CNN for Image Classification

- Each image is 6x6 pixel

Black \rightarrow 1, white \rightarrow 0

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1

CNN Filter (Kernel)

- CNN applies a filter (kernel) to the input image

Input

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (kernel)

0	0	1
0	1	0
1	0	0

Example: a 3x3 filter (9 weights);
the weights are determined by
training (i.e. **backpropagation**).

CNN Filter (Kernel)

- Overlay the filter on the image and take the dot product

Input

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (kernel)


0	0	1
0	1	0
1	0	0

dot product = $0 \times 0 + 0 \times 0 + 1 \times 1$
 $+ 0 \times 0 + 1 \times 1 + 0 \times 0$
 $+ 1 \times 1 + 0 \times 0 + 0 \times 0 = 3$

CNN Filter (Kernel)

- Add a bias term and save this as the first element of a **feature map**


$$3 \text{ (dot product)} + -2 \text{ (bias)} = 1$$



CNN Filter (Kernel)

- Slide the filter by 1 pixel for the second element in the feature map.


$$1 \text{ (dot product)} + -2 \text{ (bias)} = -1$$



CNN Filter (Kernel)

- And repeat


$$0 \text{ (dot product)} + -2 \text{ (bias)} = -2$$



CNN Filter (Kernel)


- Same for the second row of the feature map

$$1 \text{ (dot product)} + -2 \text{ (bias)} = -1$$




CNN Filter (Kernel)

- We *convolve* the filter with the input to get the feature map.




CNN Filter (Kernel)

- The filter shares the same weights and bias for all elements in the feature map. → Reduce parameters.



CNN Filter (Kernel)

- Each element of the feature map contains information of a group of neighboring pixels → takes correlations into account.




CNN Filter (Kernel)

- Feed the feature map into an activation function, say ReLU.

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

Feature map



ReLU




1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

Post-ReLU feature map

CNN Pooling


- Apply a pooling to the new feature map, e.g. **Max Pooling** which selects the maximal input



Post-ReLU feature map

CNN Pooling


- Pooling moves in a way that's *not* overlapping



Post-ReLU feature map

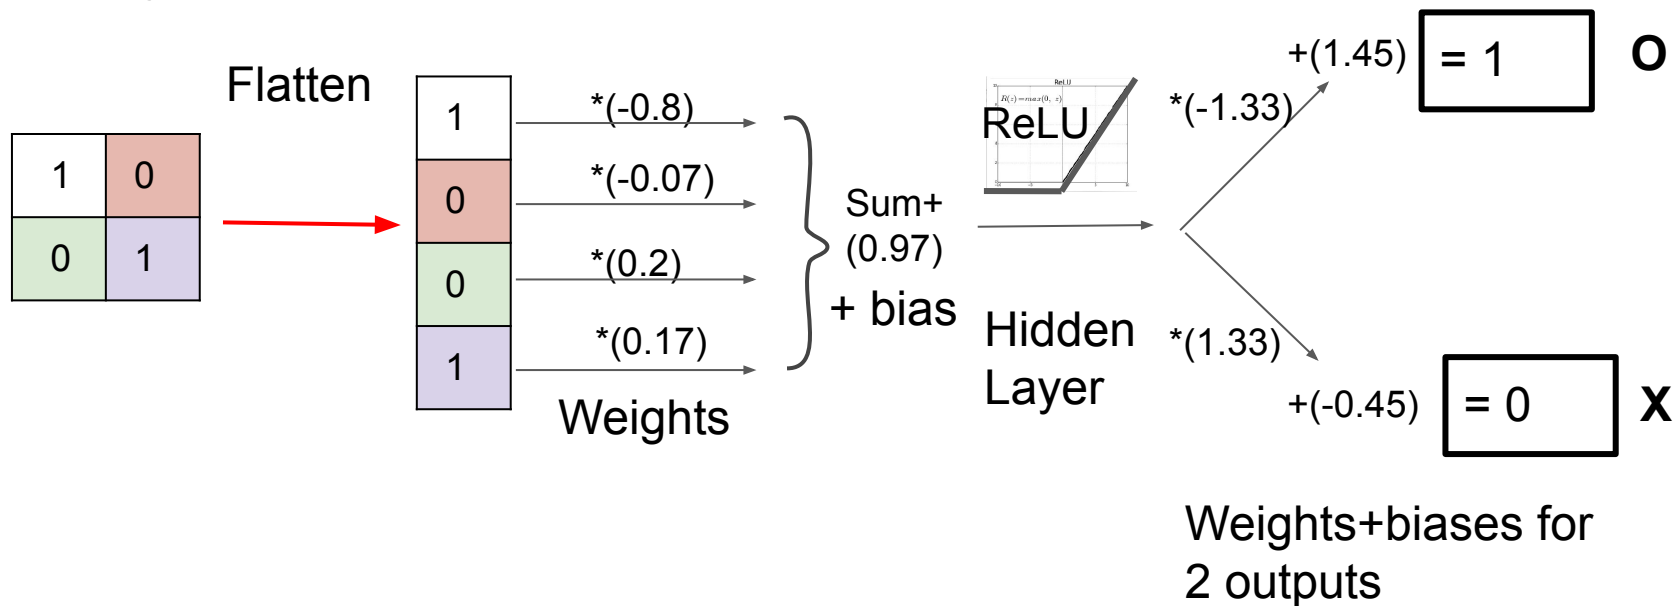
CNN Flattening and fully-connected layers

- **Flatten** the pooled layer into a 1D array and feed it into a fully-connected NN for the classification




CNN Flattening and fully-connected layers

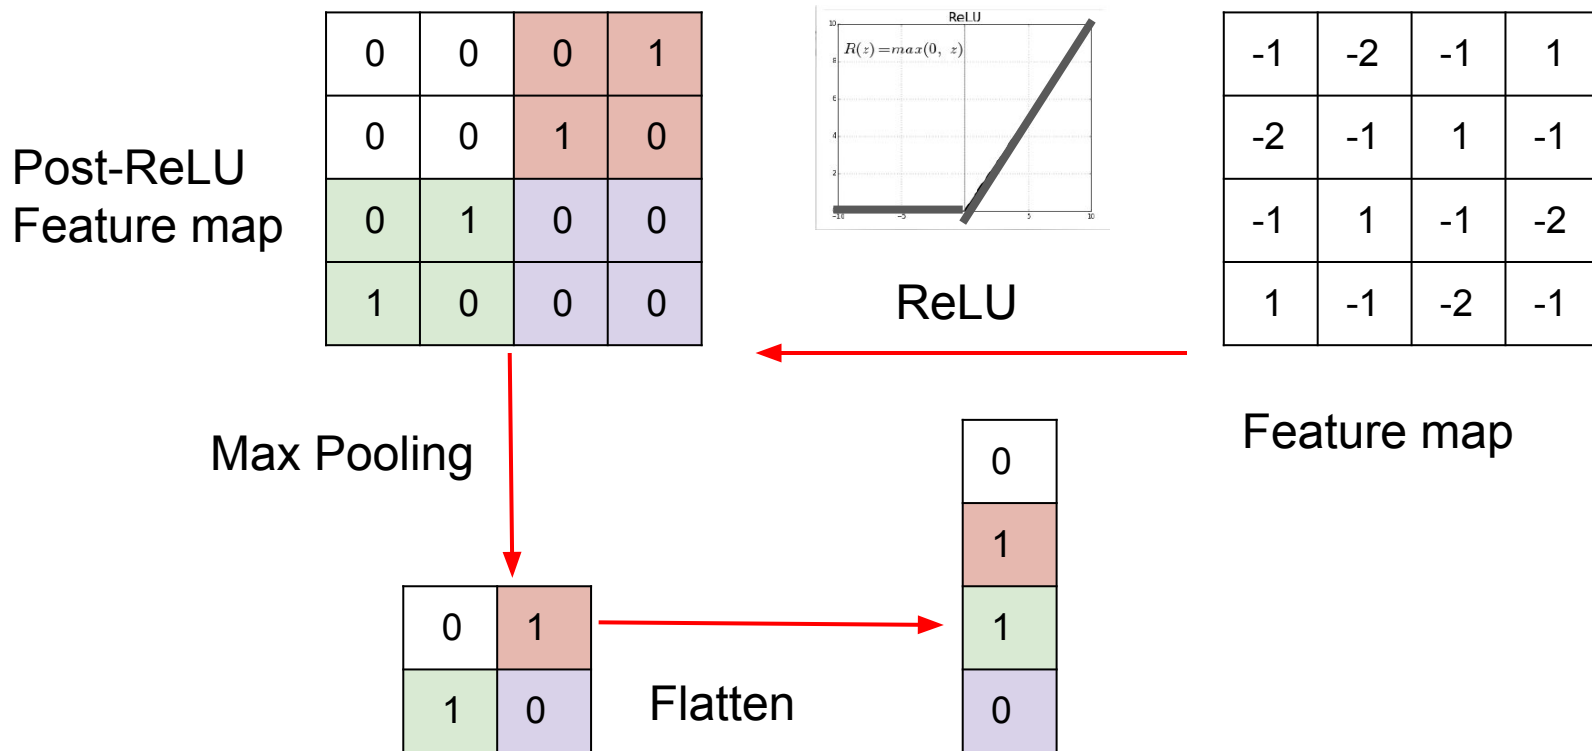
- **Flatten** the pooled layer into a 1D array and feed it into a fully-connected NN for the classification



CNN for Image Classification

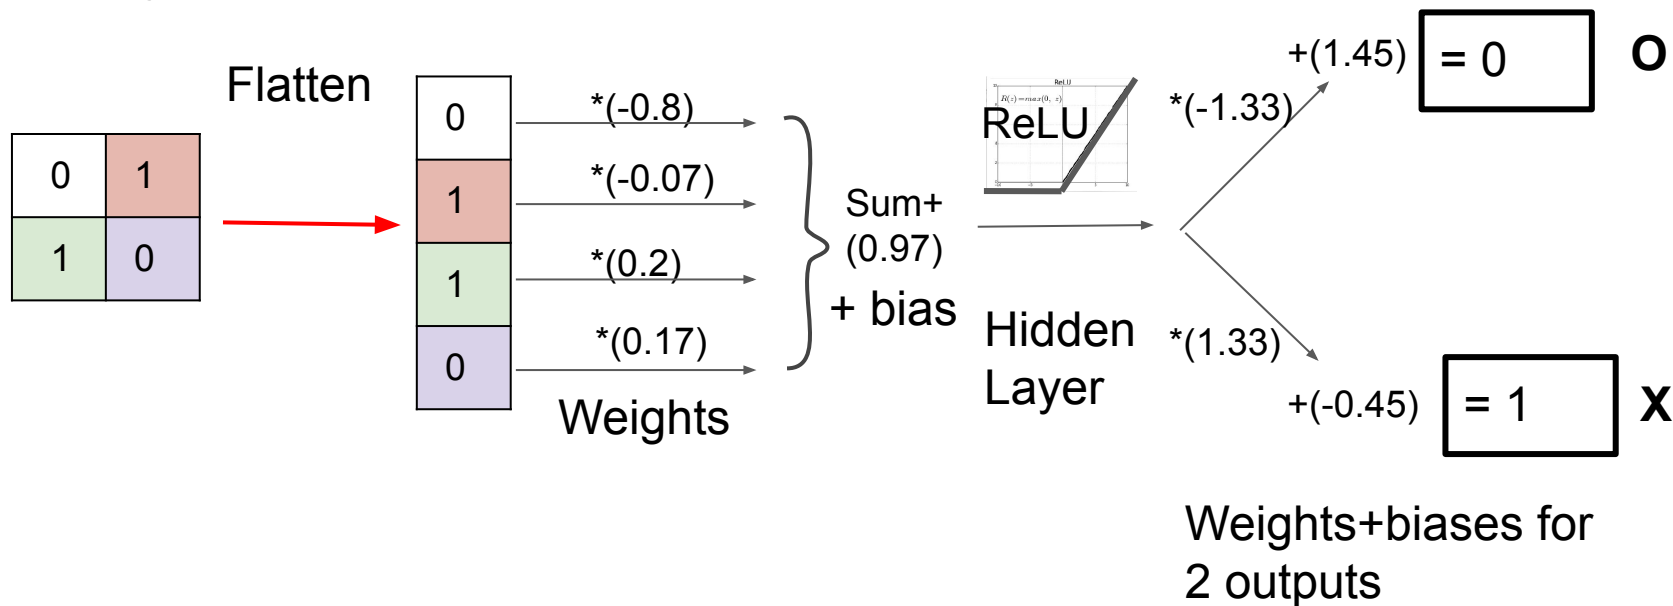


CNN for Image Classification




CNN Flattening and fully-connected layers

- **Flatten** the pooled layer into a 1D array and feed it into a fully-connected NN for the classification




How CNN works

- To see why CNN is powerful, note that the (max) pooled layer selects the parts where the filter matches the input the best.



Input image



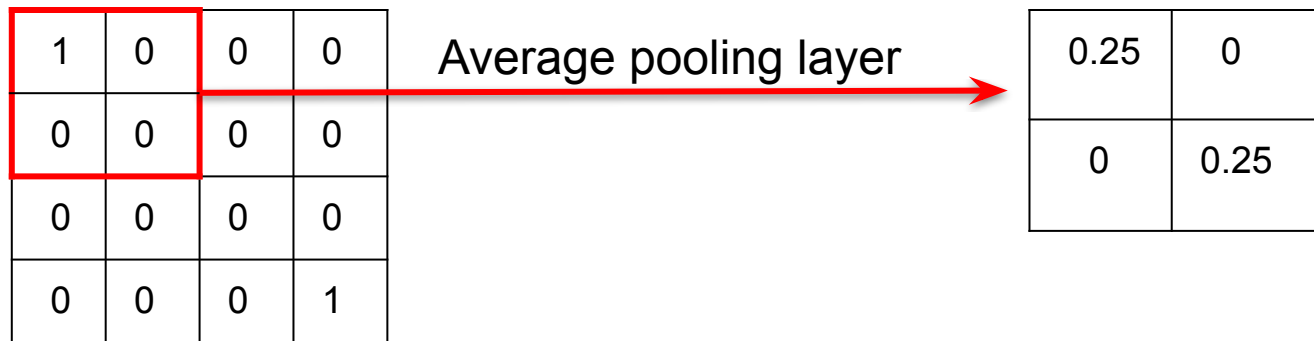
Filter

1	0
0	1

Max-pooled layer

CNN Pooling

- Another choice is **Average Pooling** which takes the average of each region



Post-ReLU feature map

CNN Strides

(a) Stride = 1

1	2	3	1	3	5
2	2	5	4	2	5
0	6	9	6	2	2
2	0	1	9	4	0
5	5	4	6	7	6
6	1	3	7	1	5

 $*$

1	0	-1
1	0	-1
1	0	-1

 $=$

-14	-1	10	-1
-11	-11	7	12
-7	-10	1	13
5	-16	-4	10

(b) Stride = 2

1	2	3	1	3	5
2	2	5	4	2	5
0	6	9	6	2	2
2	0	1	9	4	0
5	5	4	6	7	6
6	1	3	7	1	5

 $*$

1	0	-1
1	0	-1
1	0	-1

 $=$

-14	10
-7	1

CNN Padding

- We can apply padding to allow more space for the filter to cover the image and preserve the size of feature maps.

0	0	0	0	0	0	0	0
0	1	2	3	1	3	5	0
0	2	2	5	4	2	5	0
0	0	6	9	6	2	2	0
0	2	0	1	9	4	0	0
0	5	5	4	6	7	6	0
0	6	1	3	7	1	5	0
0	0	0	0	0	0	0	0

 $*$


1	0	-1
1	0	-1
1	0	-1

 $=$

-4	-5	-1	3	-5	5
-10	-14	-1	10	-1	7
-8	-11	-11	7	12	8
11	-7	-10	1	13	13
-6	5	-16	-4	10	12
-6	4	-7	-1	2	8

CNN image channel

- A color image has three channels (R/G/B) and thus needs three layers of kernels.



CNN hyperparameters

- Filter/Kernel size (height and width of the kernel)
- Strides and padding
- Data format/channel numbers
- Ways of pooling
- Other hyperparameter for the (fully-connected) NN, e.g. activation functions.

In-Class Exercise and Lab for this week

- For in-class exercise this week we'll use CNN to classify 0-9 with MNIST dataset.
- For the Lab this week, we continue with the same W/Z v.s. QCD jet dataset and use CNN for the classification task.

Backup

Outline

- Why need CNN; problem w/ deep NN


0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1

Outline

- Why need CNN; problem w/ deep NN

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0



Outline

0	0	1
0	1	0
1	0	0

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
